

Stage 2 - Processing

User guide | v1.4.0

Obsidian Compact edition. A practical guide to processing documents, checking the result, and turning verified corrections into better software.

Stage 1 gathers the documents. Stage 2 classifies, names and organises them. Stage 3 uploads the prepared folders. Completing Stage 2 does not itself upload anything.

The important distinction: processing prepares documents; the post-run accuracy audit makes suggestions; an AI review checks those suggestions against the documents; a separate learning review decides whether the software should change.

Your usual route

1. Choose the care home's **[Files]** folder and confirm the **[Processed]** destination.
2. Check Settings, start processing, and review the run estimate.
3. Let processing and any pending batch follow-up finish.
4. Let the optional post-run accuracy audit finish, then open **Reports**.
5. Use **Review audit...** to have Luna or Opus check the real documents and record decisions.
6. Use **Learn from corrections...** to have Fable or Sol assess justified software improvements.

Find the right page

Topic	Page
The new workspace and its controls	2
Folder setup and your first run	3
Live processing, batches and recovery	4
Reading progress and the post-run audit	5
Reports: audit versus review ledger	6
Launching the document review	7
Naming collisions, ranking and review quality	8
Learning from corrections and changing code	9
Settings and API usage	10
Discord and Telegram updates	11
Troubleshooting and completion checklist	12

Edition: 6 September 2026. Examples are generic. This guide contains no account credentials or worker documents.

1. The Obsidian Compact workspace

The main screen keeps the essentials visible: the selected care home, the current phase, work completed, the current document, recent activity and three headline figures. Detailed information is still available when you need it.

Control	What you use it for
Choose care-home folder	Select the folder containing worker subfolders, normally the care home's [Files] folder.
Start processing	Scan and confirm a new processing run. Read the estimate before submitting.
Check batch status	Check saved provider jobs and apply ready results. A saved batch may have more than one phase.
Stop	Request a safe stop. It is not a pause button or a promise of instant interruption.
Preview / Hide preview	Show or hide the current document without crowding the main work area.
Details & full log	Open the full activity log, detailed counters, and Folder & utilities.
Review audit...	Prepare an account-selected Luna or Opus document review.
Learn from corrections...	Prepare a Fable or Sol review of accumulated correction evidence.

The top navigation stays available

- **Jobs:** find saved jobs and their status, then check the relevant care-home run.
- **Reports:** choose the automatic audit or the AI review/correction ledger.
- **API Usage:** inspect Stage 2 and its workflow tools' recorded API usage.
- **Tools:** open the additional processing utilities available in this installation.
- **Guide:** open this guide.
- **Settings:** change processing options, limits and notification preferences.

Fewer headline numbers, not less information

Workers complete counts finished worker folders. **Needs review** counts audit flags, not proven naming errors. **Estimated run cost** is an estimate based on recorded usage and the configured exchange rate, not an account balance.

The app icon appears beside Stage 2. The top ribbon is black. On supported Windows versions, the system caption bar containing minimise, maximise and close is black too; older Windows versions may use the operating system's dark fallback.

Processing uses the Anthropic API key in Settings. The two external AI-review workflows use the Claude or Codex account you select. These are separate billing and usage contexts.

2. Folder setup and your first run

Choose the care-home folder, not an individual worker folder and not the parent containing several care homes. Files stored only in the cloud must be downloaded before processing.

```
Care Home/  
  Care Home [Files]/  
    Worker A/  
    Worker B/  
  Care Home [Processed]/
```

Before starting

1. Open **Settings** and confirm an Anthropic API key is available. Keep the existing classification model unless you intend to change processing behaviour.
2. Check the worker, file, spend and maximum-file-size limits. The estimate screen explains what will be sent, skipped or deferred.
3. For normal handover, enable **Move processed workers to a destination folder** under File movement, then choose the separate **[Processed]** destination when prompted.
4. Keep PDF conversion enabled when appropriate. LibreOffice is needed to preserve Office-document layout in full PDF conversion. Without it, Word may fall back to text-only conversion and other Office files may remain unconverted.
5. Enable **Accuracy audit after processing** under Post-run checks if you want the automatic naming check at the end. This consumes additional API usage and must fit the run budget.
6. Choose the **[Files]** folder, wait for its scan, and press **Start processing**. Review the mode, scope, warnings and estimated cost before confirming.

What happens to the documents

Stage 2 converts supported files, classifies them against its controlled names, performs its worker-level checks and organises the results. Typical finished folders contain **Overwrite Documents** and **Bulk**, with numbered bulk batches for upload. The controlled routing rules decide which documents belong where.

With Move mode enabled, a worker moves only after its processing is finalised. A worker already present at the destination is treated as already done; do not assume it will be silently replaced. With Move mode disabled, changes happen in the selected source folder.

Flattening is a utility, not a prerequisite

Details & full log > Folder & utilities > Flatten folders only brings nested files into the worker folder without calling the AI. It still changes folder organisation. Normal processing handles its own preparation, so do not flatten a pending run or repeat it merely because a run is waiting.

Keep a recoverable copy of your original source set. Do not manually move, rename or replace files belonging to a pending batch or an active review.

3. Live processing, batches and recovery

Live mode

The app processes documents while it remains open and online. You can follow the current document, preview and activity. Depending on Settings, an unknown document may prompt you to define it or be filed as a descriptive **Other** item for later checking.

Read an unknown document before defining it. A controlled compliance type belongs in the controlled vocabulary; miscellaneous material should not be turned into a new compliance type merely to avoid an unknown result.

Overnight Batch mode

Batch mode submits work to the provider for asynchronous processing. The confirmation screen separates primary classification, follow-up allowance, required finishing work and the optional audit. Use that estimate rather than assuming every part of a batch run has the same price.

1. Wait for submission to finish and the app to report a saved pending batch.
2. When ready to continue, select the same care-home folder and press **Check batch status**.
3. Stage 2 checks the saved provider jobs. If results are ready, it applies them and continues the run.
4. Some results need a follow-up batch. Follow-up work can be split into several bounded requests; a large follow-up is not automatically an error.
5. Continue checking until the run's finishing work and enabled audit have reached a terminal result. A submitted batch is not a completed worker folder.

Only close the app when it has finished local work and says the batch is safely pending. Keep it open while it applies results, organises workers or runs the audit. Provider turnaround is variable; a quiet period alone does not mean failure.

If a submission is ambiguous

An internet interruption can happen after the provider accepted a batch but before Stage 2 received its acknowledgement. Stage 2 then cannot safely assume whether that submission exists. The saved planned request and attempted submission are retained so the provider can be reconciled with the local state.

- Use **Check batch status** and follow the recovery explanation.
- Do not delete the hidden batch-state file or force a fresh live run to bypass the warning.
- Do not repeatedly submit the same work. Recovery must first identify any already accepted requests.
- If files were moved or modified since submission, restore the expected source or investigate the mismatch before resuming.

Cache and stop behaviour

Content-hash records help avoid repeating settled work. They do not mean every finishing check, audit or intentionally forced reprocess is free. **Stop** preserves known completed work where supported; a request already in flight may finish first and can still be billed.

4. Reading progress and the post-run audit

The progress display describes the **current phase**. Processing, waiting for a provider batch and checking filenames are different phases; do not interpret one phase's percentage as completion of the whole pipeline.

What the post-run accuracy audit does

When enabled and affordable within the remaining run budget, it re-examines the final eligible documents and compares their existing names with the classification evidence. Concerned cases can receive an additional opinion. It writes an audit report; it does not rename documents simply because it suggests a different name.

Display	How to read it
384 of 600 documents checked	The current audit has settled 384 document outcomes. This example is illustrative, not a live count.
64%	Work completed in that audit, not 64% naming accuracy.
Current document / current check	The document and step being worked on, including rendering or waiting for an AI response.
Elapsed waiting time	Time spent in the current wait. It explains a stationary bar without inventing progress.
About ... remaining - estimate	A timing estimate once enough useful progress exists. It can change.
Time remaining unavailable	There is not enough reliable evidence for an estimate, or the current wait makes one misleading.
Needs review	Flagged results to inspect; not a count of confirmed mistakes.

Completion, skipping and interruption

Audit complete means the report was written after the audit reached completion. Open **Reports** to inspect it. A complete report still contains AI judgements that can be wrong.

Audit skipped means it did not run, for example because it was disabled or the remaining budget was insufficient. Completed processing is not undone. Do not describe a skipped audit as passed.

Stopped or failed means the audit is incomplete. Do not submit a partial report as if every document was checked. Restarting the audit checks documents again and can consume additional API usage; the Stop button is not a pause/resume mechanism.

If the bar stops moving

Check the current step, elapsed wait, recent activity and **Details & full log**. A slow model response is different from a recorded error. Do not repeatedly click Start, cancel and retry, or launch an external correction review while the app is still processing the same documents.

5. Reports: two different kinds of evidence

Press **Reports** to answer **Which report would you like?** The automatic audit and the review ledger are intentionally separate.

1 - Post-run filename audit

Choose **Browse audit reports**, select the right care home and run, then open its report. Stage 2's current output is **Filename_Audit_Report.csv**, with a number added when needed to avoid overwriting an older report.

Read the current filename, suggested filename, full file path, review status, confidence, reason and evidence together. A confidence score is the AI's assessment, not proof. Even a high-confidence suggestion must be checked against the actual document before a correction.

CSV output can have companion **Summary**, **Orientation** and **Tables** files. Keep the companions with the main report; they replace the separate tabs of the older Excel output. Use the main audit CSV, not a summary or tables manifest, when starting an AI review. Existing Excel **.xlsx** audits remain supported by the review workflow.

2 - AI review / correction ledger

Choose **Open AI review ledger** to open **Master_Filename_Review_Ledger.xlsx**. Luna or Opus records what it actually checked, what it kept or changed, what remains uncertain, and the evidence needed for later improvement work.

The ledger is shared across accounts, providers and runs. On an existing setup it remains in the configured Stage 2 Audit Review folder, beside **review_records.jsonl**, the recovery journal. Do not create a blank replacement merely because you switch from Opus to Luna or start another care home.

Open legacy Misnaming Record appears when that compatible correction index exists. It is useful history, but it is not a replacement for the master review ledger and journal.

How the records relate

Record	Question it answers
Filename audit	What did Stage 2 suspect after processing?
AI review ledger and journal	What did the reviewer verify, decide and actually correct?
Per-request review summary and change record	What happened in this exact review, with before/after evidence and outstanding cases?
Learning review and tests	Which software changes were justified, implemented and verified?

Close a workbook in Excel before a reviewer updates it. A locked workbook must be reported and reconciled; it must not be silently replaced or treated as successfully saved.

6. Review audit... - Luna or Opus

This step independently checks the audit against the real documents. It is not another blind rename pass. The two available choices are **Luna - Codex, high effort** and **Opus - Claude Code, high effort**. These are configured review choices, not performance guarantees.

Prepare the handoff

1. Wait until Stage 2 and the exact post-run accuracy audit finish. Click **Review audit...**
2. Choose the model and a matching provider account. Use **Expected account email** when you want a specific signed-in identity checked.
3. Confirm the completed main audit, care-home name, audited document folder, original **[Files]** folder, master ledger and Stage 2 source checkout. Every review needs the current helper and naming implementation to prepare its queue and records, even without corrections.
4. Confirm **The selected post-run accuracy audit has finished**. File existence alone is not confirmation.
5. Decide whether to allow evidence-backed filename corrections. Without that permission, the review must remain a review/proposal, not apply changes.
6. Press **Verify & prepare**. The app checks account identity and model support and writes the exact request. Preparing is not starting an AI task.
7. Read the verified account and handoff details, then press **Open selected AI. Copy command** and **Open context** are available too.

What opens

An account-isolated terminal opens in the persistent audit-review workspace. Claude receives the prefilled **/stage2-review-audit** command and the exact request-file path. Codex receives the equivalent request-file prompt. Both are directed to the same role rules, naming rules, records and shared file-based memory.

The launcher does not claim to inject text into an already open VS Code extension. Existing IDE windows can retain a different account environment. The isolated terminal avoids silently using that other account; you may open the same request from a correctly signed-in IDE yourself.

What the reviewer must do

The desktop handoff includes all flagged/error rows, including lower-confidence cases. The reviewer must inspect the actual evidence, keep correct names, correct only justified cases when allowed, and defer unresolved cases honestly. It must use the supplied helper to plan and apply changes, maintain the master ledger and journal, and reconcile outputs before reporting completion.

The older manual command may use a confidence-greater-than-80 queue. Follow the scope written in this exact request, not an assumption from an older session. Neither a launch notification nor a terminal opening proves the review finished.

7. Naming collisions and a trustworthy review

Official naming and routing rules come from the Stage 2 implementation and the supplied naming-rule snapshot. The reviewer should not invent a near-match, shorten a type arbitrarily or change rules to fit one example.

Ranked families: check every peer

Some document families use their numbered suffix to express order. In those families, **the higher number is the preferred copy**. Simply adding the next number would incorrectly promote a corrected document without comparing its quality with the existing copies.

When a correction enters or changes a ranked family, the helper requires evidence and scores for **all peers in that worker's affected family**, not only the colliding file. It then plans the entire order using the actual Stage 2 rules. Readability, completeness, signing and actual document dates can matter according to the family.

Illustrative ranked order:

Document Type.pdf	lower-ranked copy
Document Type (01).pdf	next copy
Document Type (02).pdf	preferred copy

This illustration explains suffix order, not a universal instruction to create three files or to rank every document type.

Unranked families: avoid collisions without implying quality

For an unranked family, a collision may use the next available number. The reviewer must still preserve the correct controlled name, extension and routing. It must not overwrite an existing document to make its proposed name fit.

Dates, source bytes and evidence

- Use the date evidenced by the document where the naming rule calls for it, not the file's modified date or today's date.
- Do not assume a signature is present just because a filename says signed.
- The transaction helper checks the saved inventory and hashes, locks the relevant folders, creates backups and plans collision-safe changes.
- If a source changed, a peer is missing, a folder is busy or the evidence is incomplete, stop and resolve the mismatch. Do not bypass the helper with ad-hoc filesystem commands.
- A filename correction does not authorize deleting, rewriting, merging, splitting or rotating the underlying document.

What a completed review should leave

Look for a review summary, recorded decisions, applied-change evidence where applicable, a reconciled ledger/journal and a clear list of unresolved cases. "Nothing needed changing" is valid if supported. "The model exited successfully" is not sufficient verification.

8. Learn from corrections... - Fable or Sol

This is a separate reviewer with a different job: critically inspect the accumulated review evidence and decide whether a general software change is warranted. It may find that a previous correction was wrong, too specific to generalise, or not caused by a code defect.

The available choices are **Fable - Claude Code, xhigh effort** and **Sol - Codex, xhigh effort**. The launcher checks the chosen account and supported configuration. It never silently replaces an unavailable model or account.

Launch the learning review

1. Finish and reconcile the document review first. Do not send an unfinished spreadsheet as settled evidence.
2. Click **Learn from corrections...** and choose the model and provider account.
3. Confirm the existing master ledger and the actual Stage 2 **Git source checkout**, not the installed application folder.
4. Keep the same persistent shared-context root unless you deliberately want a new context location.
5. Leave code-change permission unticked for a recommendation-only review, or allow verified source-code changes with regression tests and quality checks.
6. Press **Verify & prepare**, check the handoff, then **Open selected AI**.

The AI's required sequence

- Reconcile the ledger with its journal and read prior verified lessons and unresolved cases.
- Review the evidence behind the previous reviewer's decisions; do not blindly trust every renamed file.
- Find a reproducible cause in naming, classification, ranking, routing or another in-scope function.
- Consider unaffected document types and difficult counterexamples before broadening a rule.
- Add a regression test for the defect and run relevant neighbouring tests.
- Make the smallest justified general change only when code changes are allowed.
- Record exactly what changed, what tests passed, what remains uncertain and whether further evaluation is needed. Do not mark evidence resolved just because a patch exists.

Context is shared deliberately

Opus and Luna share the **audit-review** role folder. Fable and Sol share a separate **code-learning** role folder. Each role has **MEMORY.md** plus an exact request folder containing rule snapshots and manifests. Both providers read those files; their native chat histories remain separate.

Changing accounts does not erase the ledger or role memory. Changing roles does not give the learning reviewer permission to rename care-home documents. The button does not authorize publishing a GitHub release or installing a new desktop build; those are separate, explicitly managed release steps.

9. Settings and API Usage

Processing settings affect the main document-processing engine. Choosing Luna, Opus, Fable or Sol in an external-review dialog does not change the engine's classification model.

Setting	Practical effect
Anthropic API key	Supplies processing API access. The normal installation uses the operating system credential store; leaving a replacement field blank retains the existing key.
Model / advanced models	Changes the classification model used by the processing engine. Keep an intentional, tested setting; this guide does not replace the in-app model/cost information.
USD to GBP rate	Converts recorded usage into the displayed GBP estimate.
Worker, file and spend limits	Bound a run. Review which limit was reached before changing it or continuing.
Skip files larger than (MB)	Leaves over-limit files unsent and reports them. It is separate from provider batch-payload chunking.
Image resolution / adaptive pages	Changes what page evidence is sent and when. Lower cost is not automatically equivalent evidence quality.
Skip API when clearly identified	Uses a shortcut when filename/text appears sufficient. Understand the accuracy tradeoff before enabling it.
Auto-file unrecognised as Other	Avoids live definition prompts; unresolved documents still need attention.
Local page orientation	Off, audit-only/shadow mode, or optional high-confidence correction. An orientation flag is not itself a proven rotation error.
PDF conversion	Converts supported formats; full Office conversion needs LibreOffice.
File movement	Moves finished workers to the chosen destination, or processes in place when disabled.
Accuracy audit after processing	Runs the optional report-only check, subject to remaining budget and successful completion.
Notifications...	Opens Discord and Telegram delivery settings.

Understanding API Usage

Open **API Usage** from the top navigation. The combined view includes Stage 2 processing and its integrated workflow tools, with scope controls for the included tools. Other pipeline applications have their own usage pages.

Use the tracker to inspect recorded usage over time and by the available scopes/models. The main screen's run estimate is a convenient summary, not a provider invoice, prepaid balance or proof that every external subscription session was captured.

Claude Code and Codex review sessions use the selected subscription/account. Their account quotas and session limits are separate from the processing API key. Check AI Account Manager or the provider's own account view for those limits; do not assume the Stage 2 API cost figure includes them.

10. Discord and Telegram updates

Open **Settings > Notifications...** Enable Discord, Telegram, both or neither. Notifications report meaningful lifecycle events; they do not upload reports, attach documents or start an AI review automatically.

Configure and test

1. Enable the channel you want and enter its exact destination: a Discord channel ID or a Telegram chat ID/channel.
2. Reference an existing local credential file and its token variable, or enter a new bot token for storage in the operating system credential store.
3. Leave **New bot token** blank to keep the existing credential. A saved OS credential takes precedence over environment/file references.
4. Choose whether to include aggregate progress updates and how long a wait should be before its first update.
5. Press **Save settings**, then **Test enabled channels**. Wait for the actual delivery result; "queued" is not "delivered".

An existing Claudius Messenger configuration can be referenced locally. Bot tokens must not go into source code, screenshots, public installers or this guide. Telegram requires an explicit destination; the app does not guess who should receive the messages.

What messages mean

Event	When it is reported
Run/phase starts	Preparation, processing, organising, audit and other evidenced phase changes.
Batch or follow-up submitted	Request/batch counts, with a reminder that submission is not completion.
Worker completion	Aggregate completed-worker count, without names.
Progress	Optional 25%, 50%, 75% and 100% milestones, throttled to avoid a stream of tiny updates.
Extended wait	After the configured interval, then at most every 30 minutes; waiting alone is not labelled a failure.
Attention needed / failure	A brief controlled reason; open the app for the full details.
Audit finished	Checked and flagged counts, with a reminder to inspect Reports.
AI review launched	The selected review phase started; it does not claim the external AI has finished.
Run finished	A summary that distinguishes completed, skipped, disabled or failed audit status.

Messages use aggregate information, not worker names, document text, filenames, full paths, prompts or credentials. Duplicate events are suppressed where possible.

Delivery is best-effort. A channel failure does not stop document processing or the other channel. Check the in-app delivery result when a test fails. Unsent updates can be lost when the app closes; provider retries can occasionally duplicate a message. The application records and reports remain the authoritative run evidence.

11. Troubleshooting and completion checks

What you see	What to do
A pending batch on this folder	Use Check batch status for the same saved run. Do not start a duplicate submission.
Ambiguous batch submission	Reconcile saved state with provider jobs through recovery. Keep the state files and original inputs intact.
A limit or insufficient remaining budget	Read the estimate/status and decide whether to change the relevant limit. A skipped audit is not a passed audit.
Cloud-only or missing files	Make the files available locally, then retry the appropriate unfinished work.
A file was changed or moved after submission	Restore/reconcile it before recovery or review. Do not force a stale plan to apply.
Review cannot open while Stage 2 is busy	Wait for processing, recovery and scanning to finish. The current run has not been interrupted.
Wrong account, unavailable model or unsupported effort	Open the intended account in AI Account Manager/provider, verify sign-in and select explicitly. No fallback is chosen for you.
Review requires a source checkout	Select the current Stage 2 source folder containing the transaction helper and naming implementation.
Workbook is open / records could not save	Close the workbook, then reconcile the retained journal and outputs. Do not overwrite the ledger.
No notification arrives	Check enabled channels, destination, credential reference and the actual test result. Processing may still be succeeding.

Before Stage 3 uploads

- Confirm all intended worker folders completed, including any follow-up work.
- Review skipped, unknown, unconverted and error cases in **Details & full log**.
- Confirm whether the optional accuracy audit completed; do not infer it from a full processing bar.
- Check and reconcile the external document-review result if you launched one.
- Confirm the intended processed folder and handover roster are the ones Stage 3 will use. Roster-refresh warnings need attention; folder completion alone is not proof of matching/upload readiness.

Before calling a software improvement finished

Check the learning report, evidence, tests and remaining uncertainties. A code edit or successful terminal exit does not update the installed app. Verify the build, installer, installed executable, shortcut and GitHub release separately.

Further reference

The source checkout includes **docs/ai-review/WORKFLOW_GUIDE.md**, the review/naming/learning rules, **docs/TROUBLESHOOTING.md**, **docs/NOTIFICATIONS.md** and **docs/AI_WORKFLOW_INTEGRATION.md**. Advanced vocabulary changes should start with **docs/VOCABULARY_GUIDE.md**.

Maintainers: update **docs/USER_GUIDE.md**, then run **tools/build_user_guide.py**. Release both source and PDF.