

Stage 2 - Processing

User guide | v1.5.9 Obsidian / Jade

Stage 1 gathers the documents. Stage 2 classifies, names and organises them. Stage 3 uploads the prepared folders. Completing Stage 2 does not itself upload anything.

The important distinction: processing prepares documents; the post-run accuracy audit makes suggestions; an AI review checks those suggestions against the documents; a separate learning review decides whether the software should change.

Your usual route

1. Choose **[Files]**, confirm **[Processed]**, check Settings and the run estimate.
2. Let processing, follow-up batches and the optional **Accuracy Audit** finish.
3. Open **Reports**, then **AI Document Review** to check evidence and record decisions.
4. Use **Improve Stage 2** for a separate review of justified software improvements.

Find the right page

Topic	Page
The new workspace and its controls	2
Folder setup and your first run	3
Live processing, batches and recovery	4
Ranking needs attention and targeted recovery	5
Reading progress and the post-run audit	6
Reopening a window and watching saved sessions	7
Reports: audit versus review ledger	8
Launching the document review	9
Naming collisions, ranking and review quality	10
Learning from corrections and changing code	11
Settings and API usage	12
Discord, Telegram and Slack updates	13
Troubleshooting and completion checklist	14
Reliable recovery and source exceptions	15

Edition: 15 September 2026. Examples are generic. This guide contains no account credentials or worker documents.

1. The Obsidian Compact workspace

The idle view stays clean. Once a run starts, progressive panels reveal the phase, progress, current document and activity; errors and deeper details appear when relevant.

Control	What you use it for
Choose care-home folder	Select the folder containing worker subfolders, normally the care home's [Files] folder.
Start processing	Scan and confirm a new processing run. Read the estimate before submitting.
Check batch status	Check saved provider jobs and apply ready results. A saved batch may have more than one phase.
Stop	Request a safe stop. It is not a pause button or a promise of instant interruption.
Preview / Hide preview	Show or hide the current document without crowding the main work area.
Details & full log	Open the full activity log, detailed counters, and Folder & utilities.
AI Document Review	Prepare an account-selected document review.
Improve Stage 2	Prepare a software-learning review of accumulated correction evidence.

The workspace uses Graphite surfaces, restrained Deep Jade accents, real buttons and progressive active panels. **Accuracy Audit** finds possible mistakes; **AI Document Review** checks pages; **Improve Stage 2** studies the ledger for justified software changes. Corrections never grant code authority.

The top navigation stays available

- **Jobs:** find saved jobs and their status, then check the relevant care-home run.
- **Reports:** choose the automatic audit or the AI review/correction ledger.
- **API Usage:** inspect Stage 2 and its workflow tools' recorded API usage.
- **Tools:** open the additional processing utilities available in this installation.
- **Guide:** open this guide.
- **Settings:** change processing options, limits and notification preferences.

Fewer headline numbers, not less information

Workers complete counts finished worker folders. **Needs review** counts audit flags, not proven naming errors. **Estimated run cost** is an estimate based on recorded usage and the configured exchange rate, not an account balance.

Processing uses the Anthropic API key in Settings. The two external AI-review workflows use the Claude or Codex account you select. These are separate billing and usage contexts.

2. Folder setup and your first run

Choose the care home's **[Files]** folder, not a worker folder or a parent containing several care homes. Download cloud-only files first.

```
Care Home/
  Care Home [Files]/
    Worker A/
    Worker B/
  Care Home [Processed]/
```

Before starting

1. In **Settings**, confirm the Anthropic API key, classification model, worker/file/spend limits and maximum file size.
2. For normal handover, enable **Move processed workers to a destination folder** under File movement and select the separate **[Processed]** destination.
3. Enable PDF conversion as needed. LibreOffice supports Office-to-PDF conversion; without it, Word may become text-only and other Office files may remain unconverted. Neither route guarantees fidelity.
4. Enable **Accuracy audit after processing** under Post-run checks if wanted. Its additional API usage must fit the budget.
5. Choose **[Files]**, wait for the scan, then **Start processing**. Before confirming, check the mode, scope, warnings and estimate, including skipped or deferred files.

What happens to the documents

Stage 2 converts, classifies, checks and organises documents using its controlled naming/routing rules. Finished folders normally contain **Overwrite Documents** and **Bulk**, with numbered upload batches.

Move mode moves a worker only after finalisation. An existing destination worker is treated as already done, not silently replaced. With Move mode off, the selected source folder is changed in place.

Conversion is not a fidelity check

Preserve originals and compare important source pages/sheets with their PDFs. Conversion can lose ink/signatures, images, cells or fields, or misalign answers and labels. Fixed-height source rows may already hide content; splitting PDFs can lose form-widget appearances. Successful conversion/naming does not prove completeness. Repairing one file does not prove the general converter is fixed.

Flattening is a utility, not a prerequisite

Details & full log > Folder & utilities > Flatten folders only moves nested files into the worker folder without AI calls. It changes organisation. Normal processing prepares its own files; do not flatten a pending run or repeat this utility just because it is waiting.

Keep a recoverable copy of your original source set. Do not manually move, rename or replace files belonging to a pending batch or an active review.

3. Live processing, batches and recovery

Live mode

The app processes documents while it remains open and online. You can follow the current document, preview and activity. Depending on Settings, an unknown document may prompt you to define it or be filed as a descriptive **Other** item for later checking.

Read an unknown document before defining it. A controlled compliance type belongs in the controlled vocabulary; miscellaneous material should not be turned into a new compliance type merely to avoid an unknown result.

Overnight Batch mode

Batch mode submits work to the provider for asynchronous processing. The confirmation screen separates primary classification, follow-up allowance, required finishing work and the optional audit. Use that estimate rather than assuming every part of a batch run has the same price.

1. Wait for submission to finish and the app to report a saved pending batch.
2. When ready to continue, select the same care-home folder and press **Check batch status**.
3. Stage 2 checks the saved provider jobs. If results are ready, it applies them and continues the run.
4. Some results need a follow-up batch. Follow-up work can be split into several bounded requests; a large follow-up is not automatically an error.
5. Continue checking until the run's finishing work and enabled audit have reached a terminal result. A submitted batch is not a completed worker folder.

Only close the app when it has finished local work and says the batch is safely pending. Keep it open while it applies results, organises workers or runs the audit. Provider turnaround is variable; a quiet period alone does not mean failure.

If a submission is ambiguous

An internet interruption can happen after the provider accepted a batch but before Stage 2 received its acknowledgement. Stage 2 then cannot safely assume whether that submission exists. The saved planned request and attempted submission are retained so the provider can be reconciled with the local state.

- Use **Check batch status** and follow the recovery explanation.
- Do not delete the hidden batch-state file or force a fresh live run to bypass the warning.
- Do not repeatedly submit the same work. Recovery must first identify any already accepted requests.
- If files were moved or modified since submission, restore the expected source or investigate the mismatch before resuming.

Cache and stop behaviour

Content-hash records help avoid repeating settled work. They do not mean every finishing check, audit or intentionally forced reprocess is free. **Stop** preserves known completed work where supported; a request already in flight may finish first and can still be billed.

4. Ranking needs attention and targeted recovery

Ranking needs attention means a required date, signature or quality check did not provide usable evidence. It is different from a zero score or from a document with no applicable date. Stage 2 must not guess a score, mark a contract unsigned, or promote a copy because a request failed.

Malformed signature/date answers also defer the whole family: text booleans and missing or wrong-type fields are not evidence. Real **false** values and explicit empty dates remain valid. Saved invalid answers are not automatically rebought; older answers already converted to false/empty cannot be identified retrospectively. This check validates structure, not the AI's reading accuracy.

Incomplete worker folders stay in the source location with their saved results. Completed workers may already be in **Processed**; that does not mean the captured processing scope finished. The full **Accuracy Audit** and automatic **AI Document Review** wait until the captured scope is complete and durable.

For a Batch run, the button becomes **Recover final checks** when checks are deferred. It inspects saved evidence before downloading results. If a failed or malformed finishing check is eligible for another attempt, the app asks separately before sending new paid finishing requests and shows an estimate. Declining preserves unfinished work for later. An uncertain provider outcome is not retried this way: reconcile whether it was accepted first. Do not restart primary classification to solve a finishing attention state.

Keep the same files and saved run settings. If a document is changed, added, removed or moved, the old retry confirmation is no longer valid. Resolve the mismatch rather than deleting state or forcing a fresh run. Confirmed retries apply only to eligible unresolved finishing operations; they do not authorize a repeat primary submission.

Targeted audit-row recovery

The supported explicit-row helper can prepare a deliberately selected audit row for the existing review workflow when its original report identity and current document evidence are verified. It retains the original audit/history and binds the new request to the current source and inventory. It does not silently make every earlier **Correct** audit row part of the AI review. The standard review still includes all flagged and error rows.

Reviewers must inspect actual evidence and defer unresolved cases honestly. A signature sample covering the last three pages of a long contract is not proof about every page. Offline review or regression success is not a measured live accuracy improvement.

5. Reading progress and the post-run audit

The progress display describes the **current phase**. Processing, waiting for a provider batch and checking filenames are different phases; do not interpret one phase's percentage as completion of the whole pipeline.

What the post-run accuracy audit does

When enabled and affordable within the remaining run budget, it re-examines the final eligible documents and compares their existing names with the classification evidence. Concerned cases can receive an additional opinion. It writes an audit report; it does not rename documents simply because it suggests a different name.

Display	How to read it
384 of 600 documents checked	The audit has recorded 384 outcomes, which can include errors. This is an illustrative count, not proof that every file was readable.
64%	Work completed in that audit, not 64% naming accuracy.
Current document / current check	The document and step being worked on, including rendering or waiting for an AI response.
Elapsed waiting time	Time spent in the current wait. It explains a stationary bar without inventing progress.
About ... remaining - estimate	A timing estimate once enough useful progress exists. It can change.
Time remaining unavailable	There is not enough reliable evidence for an estimate, or the current wait makes one misleading.
Needs review	Flagged results to inspect; not a count of confirmed mistakes.

Completion, skipping and interruption

For the Employee Handbook finishing path, completed SHA-bound evidence lets the audit mark **UnableToDetermine** at 40 or below when the typed `legible` flag is true; the finishing `complete` flag is not required, because a wrong-type document is reported as incomplete for the claimed family. It does not launch an audit, rename documents, or make an extra API call; configured scope controls intake. Select all-flags to include confidence-0 results; legacy >80-only excludes them. Processing and ranking are unchanged; empty-worker and campaign safeguards are inherited. This is a review-only signal. A genuine but incomplete handbook may also be flagged. It does not establish classification accuracy or automatically correct the document.

Audit complete means the audit finished its attempts and wrote the report. It can still contain errors, unreadable inputs or incorrect AI judgements. Open **Reports** and check statuses and Notes; do not treat completion as a passed accuracy test. The naming audit examines resulting documents, not a complete source-to-output fidelity comparison.

Audit skipped means it did not run, for example because it was disabled or the remaining budget was insufficient. Completed processing is not undone. Do not describe a skipped audit as passed.

Stopped or failed means the audit is incomplete. Do not submit a partial report as if every document was checked. Restarting the audit checks documents again and can consume additional API usage; the Stop button is not a pause/resume mechanism.

Reopening a window without restarting processing

If the bar stops moving, check the current step, elapsed wait and **Details & full log**. A slow model response is not a recorded failure. Do not repeatedly click Start or begin a correction review while documents are still processing.

Open **Jobs > Running & saved sessions**. The app also opens this read-only view when it discovers an active saved operation at startup. It shows the original process, source folder, phase, counts, cost estimate and last heartbeat. It does not submit, apply, stop or take ownership of that process. Choose a saved session to watch it; reopening the viewer does not create another paid run.

Running means the recorded owner identity and recent heartbeat agree. **Stale-unverified** means the evidence is old or the owner cannot be verified; it is not proof of failure. **Owner-ended** means the recorded owner has ended or its PID was reused. **Operation-ended** only means that local operation ended: provider work, applying results, the accuracy audit or AI review may still remain. **Failed** reports a recorded operation failure. Read the outcome and normal batch status before deciding what remains; none of these labels proves accuracy.

When you close a busy dashboard, **Minimise and keep running** is the default. The process stays open on the taskbar. **Stop safely and close** requests the existing safe stop, then closes when idle; a stopped audit is not resumable. **Cancel** keeps the window open. This release does not turn a GUI-owned thread into a separate background service: ending the process, signing out, restarting or shutting down still interrupts it. Sleeping pauses local work.

Saved progress is stored under %LOCALAPPDATA%/Lifted/Stage2/runs, never in the worker folders. Diagnostics under %LOCALAPPDATA%/Lifted/Stage2Diagnostics record window lifecycle and Python/Tk/thread failures without copying document content. They help diagnose future unexpected exits; no specific historical crash cause has been established. Runs already started by older versions cannot acquire the new progress publisher retrospectively. Do not delete native batch state or writer locks to force a second run.

6. Reports: two different kinds of evidence

Press **Reports** to answer **Which report would you like?** The automatic audit and the review ledger are intentionally separate.

1 - Post-run filename audit

Choose **Browse audit reports**, select the right care home and run, then open its report. Stage 2's current output is **Filename_Audit_Report.csv**, with a number added when needed to avoid overwriting an older report.

Read the current filename, suggested filename, full file path, review status, confidence, reason and evidence together. A confidence score is the AI's assessment, not proof. Even a high-confidence suggestion must be checked against the actual document before a correction.

CSV output can have companion **Summary**, **Orientation** and **Tables** files. Keep the companions with the main report; they replace the separate tabs of the older Excel output. Use the main audit CSV, not a summary or tables manifest, when starting an AI review. Existing Excel **.xlsx** audits remain supported by the review workflow.

2 - AI review / correction ledger

Choose **Open AI review ledger** to open **Master_Filename_Review_Ledger.xlsx**. The selected document-review model records what it actually checked, what it kept or changed, what remains uncertain, and the evidence needed for later improvement work.

The ledger is shared across accounts, providers and runs. On an existing setup it remains in the configured Stage 2 Audit Review folder, beside **review_records.jsonl**, the recovery journal. Do not create a blank replacement merely because you switch from Opus to Luna or start another care home.

Open legacy Misnaming Record appears when that compatible correction index exists. It is useful history, but it is not a replacement for the master review ledger and journal.

How the records relate

Record	Question it answers
Filename audit	What did Stage 2 suspect after processing?
AI review ledger and journal	What did the reviewer verify, decide and actually correct?
Per-request review summary and change record	What happened in this exact review, with before/after evidence and outstanding cases?
Learning review and tests	Which software changes were justified, implemented and verified?

Close a workbook in Excel before a reviewer updates it. A locked workbook must be reported and reconciled; it must not be silently replaced or treated as successfully saved.

7. AI Document Review

Before Start, configure **AI Document Review** as **Model** → **Account** → **Effort**; the provider is inferred. For new settings, the default is **Opus 5 - Claude Code, High**, with **Apply supported filename corrections** on and automatic review after a complete Accuracy Audit. This step then checks the audit against real documents; it is not another blind rename pass.

Prepare the handoff

1. Wait until Stage 2 and the exact Accuracy Audit finish. Click **AI Document Review**.
2. Choose the model and a matching provider account. Use **Expected account email** when you want a specific signed-in identity checked.
3. Confirm the completed main audit, care-home name, audited document folder, original **[Files]** folder, master ledger and Stage 2 source checkout. Every review needs the current helper and naming implementation to prepare its queue and records, even without corrections.
4. Confirm **The selected post-run accuracy audit has finished**. File existence alone is not confirmation.
5. Decide whether to allow evidence-backed filename corrections. Without that permission, the review must remain a review/proposal, not apply changes.
6. Press **Verify & prepare**. The app checks account identity and model support and writes the exact request. Preparing is not starting an AI task.
7. Read the verified account and handoff details, then press **Open selected AI. Copy command** and **Open context** are available too.

What opens

An account-isolated terminal opens in the persistent audit-review workspace. Claude receives the prefilled **/stage2-review-audit** command and the exact request-file path. Codex receives the equivalent request-file prompt. Both are directed to the same role rules, naming rules, records and shared file-based memory.

The launcher does not claim to inject text into an already open VS Code extension. Existing IDE windows can retain a different account environment. The isolated terminal avoids silently using that other account; you may open the same request from a correctly signed-in IDE yourself.

At Start, the selected **Model** → **Account** → **Effort** (provider inferred read-only: Sol/Terra/Luna/Astra are Codex; Fable/Opus are Claude), scope and authority are snapshotted and fixed on resume. Existing saved selections are preserved; automatic review submits once only after a matching complete audit receipt and released locks. Pending/failed audits do not launch and the newest report is never guessed. A terminal/log viewer shows live output, not private reasoning or interactive chat; **View AI session** reopens the existing stream. Sign-in may require the user; no automatic OAuth, VS Code is optional, and CLI billing is separate from processing API billing.

What the reviewer must do

The desktop handoff includes all flagged/error rows, including lower-confidence cases. The reviewer must inspect the actual evidence, keep correct names, correct only justified cases when allowed, and defer unresolved cases honestly. Encrypted, blank or incomplete evidence must not receive invented scores, dates, signature claims or a successful-review label. A file's name or a confident earlier judgement is not proof it was readable. Use the supplied helper to plan/apply changes, maintain the ledger/journal and reconcile outputs.

The older manual command may use a confidence-greater-than-80 queue. Follow the scope written in this exact request, not an assumption from an older session. Neither a launch notification nor a terminal opening proves the review finished.

8. Naming collisions and a trustworthy review

Official naming and routing rules come from the Stage 2 implementation and the supplied naming-rule snapshot. The reviewer should not invent a near-match, shorten a type arbitrarily or change rules to fit one example.

Ranked families: check every peer

Some document families use their numbered suffix to express order. In those families, **the higher number is the preferred copy**. Simply adding the next number would incorrectly promote a corrected document without comparing its quality with the existing copies.

When a correction enters or changes a ranked family, the helper requires evidence and scores for **all peers in that worker's affected family**, not only the colliding file. It then plans the entire order using the actual Stage 2 rules. Readability, completeness, signing and actual document dates can matter according to the family.

Illustrative ranked order:

Document Type.pdf	lower-ranked copy
Document Type (01).pdf	next copy
Document Type (02).pdf	preferred copy

This illustration explains suffix order, not a universal instruction to create three files or to rank every document type.

Unranked families: avoid collisions without implying quality

For an unranked family, a collision may use the next available number. The reviewer must still preserve the correct controlled name, extension and routing. It must not overwrite an existing document to make its proposed name fit.

Dates, source bytes and evidence

- Use the date evidenced by the document where the naming rule calls for it, not the file's modified date or today's date.
- Do not assume a signature is present just because a filename says signed.
- The transaction helper checks the saved inventory and hashes, locks the relevant folders, creates backups and plans collision-safe changes.
- If a source changed, a peer is missing, a folder is busy or the evidence is incomplete, stop and resolve the mismatch. Do not bypass the helper with ad-hoc filesystem commands.
- A filename correction does not authorize deleting, rewriting, merging, splitting or rotating the underlying document.

What a completed review should leave

Look for a review summary, recorded decisions, applied-change evidence where applicable, a reconciled ledger/journal and a clear list of unresolved cases. "Nothing needed changing" is valid if supported. "The model exited successfully" is not sufficient verification.

9. Improve Stage 2

This is a separate reviewer with a different job: critically inspect the accumulated review evidence and decide whether a general software change is warranted. It may find that a previous correction was wrong, too specific to generalise, or not caused by a code defect.

This separate software-learning role recommends **Table 5.1 - Claude Code, High** for codebase-wide investigation or **Sol - Codex, High** for routine bounded learning; **Astra - Codex, Medium** is a provisional complex-case alternative. The launcher checks the chosen account and supported configuration. It never silently replaces an unavailable model or account, and filename corrections never authorize automatic code changes.

Launch the learning review

1. Finish and reconcile the document review first. Do not send an unfinished spreadsheet as settled evidence.
2. Click **Improve Stage 2** and choose Model → Account → Effort.
3. Confirm the existing master ledger and the actual Stage 2 **Git source checkout**, not the installed application folder.
4. Keep the same persistent shared-context root unless you deliberately want a new context location.
5. Leave code-change permission unticked for a recommendation-only review, or allow verified source-code changes with regression tests and quality checks.
6. Press **Verify & prepare**, check the handoff, then **Open selected AI**.

The AI's required sequence

- Reconcile the ledger with its journal and read prior verified lessons and unresolved cases.
- Review the evidence behind the previous reviewer's decisions; do not blindly trust every renamed file.
- Find a reproducible cause in naming, classification, ranking, routing or another in-scope function.
- Consider unaffected document types and difficult counterexamples before broadening a rule.
- Add a regression test for the defect and run relevant neighbouring tests.
- Make the smallest justified general change only when code changes are allowed.
- Record exactly what changed, what tests passed, what remains uncertain and whether further evaluation is needed. Do not mark evidence resolved just because a patch exists.

Context is shared deliberately

The selected document-review models share the **audit-review** role folder. Fable/Sol (and any explicitly selected Astra alternative) use the separate **code-learning** role folder. Each role has **MEMORY.md** plus an exact request folder containing rule snapshots and manifests. Providers read those files; their native chat histories remain separate.

Changing accounts does not erase the ledger or role memory. Changing roles does not give the learning reviewer permission to rename care-home documents. The button does not authorize publishing a GitHub release or installing a new desktop build; those are separate, explicitly managed release steps.

10. Settings and API Usage

Processing settings affect the main engine; review-model choices do not change classification. Theme choices are **A · Deep jade**, **B · Graphite**, and **C · Graphite + jade** (default). Preview changes live; **Save** applies and **Cancel** restores the previous settings.

Setting	Practical effect
Anthropic API key	Supplies processing API access. The normal installation uses the operating system credential store; leaving a replacement field blank retains the existing key.
Model / advanced models	Changes the classification model used by the processing engine. Keep an intentional, tested setting; this guide does not replace the in-app model/cost information.
USD to GBP rate	Converts recorded usage into the displayed GBP estimate.
Worker, file and spend limits	Bound a run. Review which limit was reached before changing it or continuing.
Skip files larger than (MB)	Leaves over-limit files unsent and reports them. It is separate from provider batch-payload chunking.
Image resolution / adaptive pages	Changes what page evidence is sent and when. Lower cost is not automatically equivalent evidence quality.
Skip API when clearly identified	Uses a shortcut when filename/text appears sufficient. Understand the accuracy tradeoff before enabling it.
Auto-file unrecognised as Other	Avoids live definition prompts; unresolved documents still need attention.
Local page orientation	Off, audit-only/shadow mode, or optional high-confidence correction. An orientation flag is not itself a proven rotation error.
PDF conversion	Converts supported formats; LibreOffice enables supported full Office conversion, but does not guarantee content/layout fidelity.
File movement	Moves finished workers to the chosen destination, or processes in place when disabled.
Accuracy audit after processing	Runs the optional report-only check, subject to remaining budget and successful completion.
Notifications...	Opens Discord, Telegram and Slack delivery settings.

Understanding API Usage

Open **API Usage** from the top navigation to inspect Stage 2 and included workflow-tool usage. Other pipeline applications have separate usage pages.

Use the tracker by scope/model and over time. The run estimate is a summary, not an invoice, prepaid balance or record of external subscription sessions.

11. Discord, Telegram and Slack updates

Open **Settings > Notifications...** Enable any combination of Discord, Telegram and Slack, or none. Updates report lifecycle events; they do not upload documents or start an AI review.

Configure and test

1. For Discord or Telegram, enter the exact channel/chat ID and a bot credential.
2. For Slack, create an app at **api.slack.com/apps**, enable **Incoming Webhooks**, and add a webhook to the exact workspace/channel. If approval is required, request it from an app administrator first; a pending request is not connected.
3. Paste Slack's URL into **New webhook URL**, or reference **SLACK_WEBHOOK_URL** in an existing local credential file. Slack's label is for reference only: the webhook's channel is chosen in Slack.
4. New credentials use the OS credential store. Leave the masked field blank to keep the saved credential, which takes precedence over environment/file references.
5. Enable your channels, choose progress/wait preferences, save, then test. Wait for **delivered**, not just **queued**, and check the intended channel.

Existing Claudius Messenger credentials can be referenced locally. Never share tokens or webhook URLs in chat, screenshots or public builds. Slack needs no separate bot server.

What messages mean

Event	When it is reported
Run/phase starts	Preparation, processing, organising, audit and other evidenced phase changes.
Batch or follow-up submitted	Request/batch counts, with a reminder that submission is not completion.
Worker completion	Aggregate completed-worker count, without names.
Progress	Optional 25%, 50%, 75% and 100% milestones, throttled to avoid a stream of tiny updates.
Extended wait	After the configured interval, then at most every 30 minutes; waiting alone is not labelled a failure.
Attention needed / failure	A brief controlled reason; open the app for the full details.
Audit finished	Checked and flagged counts, with a reminder to inspect Reports.
AI review launched	The selected review phase started; it does not claim the external AI has finished.
Run finished	A summary that distinguishes completed, skipped, disabled or failed audit status.

Messages use aggregate information, not worker names, document text, filenames, full paths, prompts or credentials. Duplicate events are suppressed where possible.

Delivery is best-effort. A channel failure does not stop processing or other channels. Closing the app can lose unsent updates; retries can duplicate messages. Application records and reports remain authoritative.

12. Troubleshooting and completion checks

What you see	What to do
A pending batch on this folder	Check batch status for this saved run; do not submit it again.
Ambiguous batch submission	Reconcile saved state with provider jobs through recovery. Keep the state files and original inputs intact.
A limit or insufficient remaining budget	Read the estimate/status and decide whether to change the relevant limit. A skipped audit is not a passed audit.
Cloud-only or missing files	Make the files available locally, then retry the appropriate unfinished work.
A file was changed or moved after submission	Restore/reconcile it before recovery or review. Do not force a stale plan to apply.
A writer-lock file remains	Windows removes it after release when possible. Open handles, permissions or sync may prevent cleanup. Check run status; do not delete it. Other lock types can be permanent.
Local orientation state could not save	Stop reading active JSON; monitor app progress. Retry when idle. Brief Windows contention is retried; persistent failure preserves the previous state. See troubleshooting.
Review cannot open while Stage 2 is busy	Wait for processing, recovery and scanning to finish. The current run has not been interrupted.
Wrong account, unavailable model or unsupported effort	Open the intended account in AI Account Manager/provider, verify sign-in and select explicitly. No fallback is chosen for you.
Review requires a source checkout	Select current Stage 2 source: the transaction helper and naming code.
Workbook is open / records could not save	Close the workbook, then reconcile the retained journal and outputs. Do not overwrite the ledger.
No notification arrives	Check enabled channels, destination, credential reference and the actual test result. Processing may still be succeeding.

Before Stage 3 uploads

- Confirm all workers and follow-up work completed. Inspect skipped, unknown, unconverted and error cases in **Details & full log**.
- Compare source/output content and preserve originals. Renaming cannot fix conversion damage.
- Verify actual accuracy-audit completion and reconcile any external document review.
- Confirm Stage 3's processed folder and roster; resolve roster-refresh warnings.

Before calling a software improvement finished

Check the learning evidence and tests. Verify the build, installer, installed executable, shortcut and GitHub release separately; code edits do not update the app.

Further reference: [docs/ai-review/WORKFLOW_GUIDE.md](#), [docs/TROUBLESHOOTING.md](#), [docs/NOTIFICATIONS.md](#), [docs/AI_WORKFLOW_INTEGRATION.md](#) and [docs/VOCABULARY_GUIDE.md](#). Rebuild guide edits with [tools/build_user_guide.py](#) and check every page. Release source and PDF; see [docs/INSTALL.md](#).

13. Reliable recovery and source exceptions

Retrieving saved batch results

Retrieving saved batch results downloads existing answers, not new document submissions. Its counter measures verified result batches, not completed workers. Temporary connection failures receive a bounded retry. Validated complete batches are cached locally so an interrupted download or later finishing retry can reuse them. Incomplete, wrong-identity or corrupted cached data is not trusted.

If downloading stops, keep the saved state and use **Check batch status** again. Do not submit the whole run again. A cache is an aid to recovery, not a replacement for the saved request history or original documents.

Exact duplicates and interrupted final checks

Stage 2 records its exact duplicate-removal plan before removing redundant copies. It retains original classification history separately and checkpoints the surviving documents before final checks. Recovery can finish an interrupted recorded removal without treating those duplicates as lost documents. Missing unique content or unexplained changes still block recovery. Old runs without this journal are not automatically rewritten to pretend an unexplained mismatch is safe.

Source exceptions are not completed workers

No documents supplied identifies an empty input folder. It is not proof that the worker is compliant or processed. Supply the missing originals, or explicitly resolve the intentionally empty folder outside the completed-document count.

Unreadable source document means Stage 2 could not obtain usable evidence. Its original identity is retained; other readable requests can continue. The worker stays incomplete. **View source issues** reopens a source-only outcome without an API key or new request. Obtain a complete original; do not overwrite an active or pending source. Replacements need a deliberate new/reconciled processing scope.

Network errors: what the app can and cannot fix

Upload diagnostics now separate sending, HTTP-status access and response reading, with numeric Windows causes where available. A missing upload acknowledgement still needs provider reconciliation before another paid submission. The app cannot guarantee an uninterrupted internet connection or infer acceptance from an error. Failed final checks use the separate costed recovery confirmation on page 5.

No model, image resolution or naming/ranking policy was reduced by these changes. Computer sleep still pauses local processing; keep the computer awake and online for unattended local work. Provider-side batch work may continue while it sleeps.